

# B561: Advanced Database Concepts

## Assignment 6

### Spring 2022

Due: Friday, April 13th, 11:59pm EST

This assignment covers:

- Object-relational databases
- Nested relational and semi-structured databases

To turn in your assignment, you will need to upload to Canvas a single file with name `assignment6.sql` which contains the necessary SQL statements that solve the problems in this assignment. The `assignment6.sql` file must be so that the AI's can run it in their PostgreSQL environment. You should use the script file to construct the `assignment6.sql` file. (Note that the data to be used for this assignment is included in this file.) In addition, you will need to upload a separate `assignment6.txt` file that contains the results of running your queries.

# 1 Formulating Query in Object-Relational SQL

For the problems in the section, you will need to use the polymorphically defined functions and predicates that are defined in the document

`SetOperationsAndPredicates.sql`

## Functions

<code>set union(A,B)</code>	$A \cup B$
<code>set intersection(A,B)</code>	$A \cap B$
<code>set difference(A,B)</code>	$A \setminus B$
<code>add element(x,A)</code>	$x \cup A$
<code>remove element(x,A)</code>	$A - x$
<code>make singleton(x)</code>	$\{x\}$
<code>choose element(A)</code>	choose some element from A
<code>bag union(A,B)</code>	the bag union of A and B
<code>bag to set(A)</code>	coerce the bag A to the corresponding set

## Predicates

<code>is_in(x,A)</code>	$x \in A$
<code>is_not_in(x,A)</code>	$x \notin A$
<code>is_empty(A)</code>	$A = \emptyset$
<code>is_not_emptyset(A)</code>	$A \neq \emptyset$
<code>subset(A,B)</code>	$A \subseteq B$
<code>superset(A,B)</code>	$A \supseteq B$
<code>equal(A,B)</code>	$A = B$
<code>overlap(A,B)</code>	$A \cap B \neq \emptyset$
<code>disjoint(A,B)</code>	$A \cap B = \emptyset$

We now turn to the problems in this section. You will need use the data provided for the `Person`, `Company`, `companyLocation`, `worksFor`, `jobSkill`, `personSkill`, and `Knows` relations. But before turning to the problems, we will introduce various object-relational views defined over these relations:

1. The view `companyHasEmployees(cname,employees)` which associates with each company, identified by a `cname`, the set of pids of persons who work for that company.  
`create or replace view companyHasEmployees as select cname,  
array(select pid from worksfor w where w.cname = c.cname order  
by 1) as employees from company c order by 1;`
2. The view `cityHasCompanies(city,companies)` which associates with each city the set of cnames of companies that are located in that city..

```
create or replace view cityHasCompanies as select city, array_agg(cname
order by 1) as companies from companyLocation group by city
order by 1;
```

3. The view `companyHasLocations(cname,locations)` which associates with each company, identified by a `cname`, the set of cities in which that company is located.

```
create or replace view companyHasLocations as select cname,
array(select city from companyLocation cc where c.cname = cc.cname
order by 1) as locations from company c order by 1;
```

4. The view `knowsPersons(pid,persons)` which associates with each person, identified by a `pid`, the set of pids of persons he or she knows.  

```
create or replace view knowsPersons as select p.pid, array(select
k.pid2 from knows k where k.pid1 = p.pid order by pid2) as persons
from person p order by 1;
```

5. The view `isKnownByPersons(pid,persons)` which associates with each person, identified by a `pid`, the set of pids of persons who know that person. Observe that there may be persons who are not known by any one.  

```
create or replace view isKnownByPersons as select distinct p.pid,
array(select k.pid1 from knows k where k.pid2 = p.pid) as persons
from person p order by 1;
```

6. The view `personHasSkills(pid,skills)` which associates with each person, identified by a `pid`, his or her set of job skills.

```
create or replace view personHasSkills as select distinct p.pid,
array(select s.skill from personSkill s where s.pid = p.pid
order by 1) as skills from person p order by 1;
```

view `skillOfPersons(skills,persons)` which associates with each job skill the set of pids of persons who have that job skill.

```
create or replace view skillOfPersons as select js.skill, array(select
ps.pid from personSkill ps where ps.skill = js.skill order by
pid) as persons from jobSkill js order by skill;
```

In the problems in this section, you are asked to formulate queries in object-relational SQL. You should use the set operations and set predicates

defined in the document `SetOperationsAndPredicates.sql`, the relations

Person  
Company  
Skill  
worksFor

and the views

companyHasEmployees  
cityHasCompanies  
companyHasLocations  
knowsPersons  
isKnownByPersons  
personHasSkills  
skillOfPersons

However, you are not permitted to use the *Knows*, *companyLocation*, and *personSkill* relations in the object-relation SQL formulation of the queries. Observe that you actually don't need these relations since they are encapsulated in these views.

Before listing the queries that you are asked to formulate, we present some examples of queries that are formulated in object-relational SQL using the assumptions stated in the previous paragraph. Your solutions need to be in the style of these examples. The goal is to maximize the utilization of the functions and predicates defined in document `SetOperationsAndPredicates.sql`.

1. **Example 1:** Consider the query "Find the pid of each person who knows a person who has a salary greater than 55000."  
`select distinct pk.pid from knowsPersons pk, worksfor w where is in(w.pid, pk.persons) and w.salary > 55000 order by 1;`

*Note that the following formulation for this query is not allowed since it uses the relation **Knows** which is not permitted.* `select distinct k.pid1 from knows k, worksfor w where k.pid2 = w.pid and w.salary >`

55000;

*Note that the following formulation for this query is not allowed since it uses the relation **Knows** which is not permitted.*

select distinct k.pid1 from knows k, worksfor w where k.pid2 = w.pid  
and w.salary > 55000;

2. **Example 2:** Consider the query "Find the pid and name of each person along with the set of his or her skills that are not among the skills of persons who work for 'Netflix' "

```
select p.pid, p.pname, set_difference((select ps.skills from
personHasSkills ps where ps.pid = p.pid), array(select unnest(ps.skills)
from personHasSkills ps where is_in(ps.pid, (select employees
from companyHasEmployees where cname = 'Netflix')))) from person
p;
```

1. Formulate the following queries in object-relational SQL.

Find the pid and name of each person  $p$  along with the set of pids of persons who (1) know  $p$  and (2) who have the AI skill but not the Programming skill and Networks.

2. Formulate the following queries in object-relational SQL.

Find each  $(c, p)$  pair where  $c$  is the cname of a company and  $p$  is the pid of a person who works for that company and who is known by everyone who works for that company.

3. Formulate the following queries in object-relational SQL.

Find the pid and name of each person who has all the skills of the combined set of job skills of the lowest paid persons who work for Google.

4. Find the following set of sets

$$\{S \mid S \subseteq \text{Skill} \wedge |S| \leq 2\}.$$

I.e., this is the set consisting of each set of job skills whose size (cardinality) is at most 2.

5. Let

$$\mathcal{S} = \{S \mid S \subseteq \text{Skill} \wedge |S| \leq 3\}.$$

Find the following set of sets

$$\{X \mid X \subseteq \mathcal{S} \wedge |X| \leq 2\}.$$

6. Let  $A$  and  $B$  be sets such that  $A \cup B \neq \emptyset$ . The *Jaccard index*  $J(A, B)$  is defined as the quantity

$$\frac{|A \cap B|}{|A \cup B|}.$$

The Jaccard index is a frequently used measure to determine the similarity between two sets. Note that if  $A \cap B = \emptyset$  then  $J(A, B) = 0$ , and if  $A = B$  then  $J(A, B) = 1$ .

Let  $t$  be a number called a *threshold*. We assume that  $t$  is a `float` in the range  $[0, 1]$ .

Write a function `JaccardSimilar(t float)` that returns the set of unordered pairs  $\{s_1, s_2\}$  of different skills such that the set of persons who have skill  $s_1$  and the set of persons who have skill  $s_2$  have a Jaccard index of at least  $t$ .

Test your function `JaccardSimilar` for the following values for  $t$ : 0, 0.25, 0.5, 0.75, and 1.

## 2 Nested Relations and Semi-structured databases

Consider the lecture on Nested and Semi-structured Data models. In that lecture, we considered the `studentGrades` nested relation and we constructed it using a PostgreSQL query starting from the `Enroll` relation.

7. Write a PostgreSQL view `courseGrades` that creates the nested relation of type

$$(cno, gradeInfo\{(grade, students\{(sid)\})\})$$

This view should compute for each course, the grade information of the students enrolled in this course. In particular, for each course and for each grade, this relation stores in a set the students who obtained that grade in that course.

8. Starting from the `courseGrades` view in Problem (a) solve the following queries:  
Find each `cno` `c` where `c` is a course in which all students received the same grade.
9. Write a PostgreSQL view `jcouseGrades` that creates a semi-structured database which stores jsonb objects whose structure conforms with the structure of tuples as described for the `courseGrades` in Problem (a). Test your view.
10. Starting from the `jcouseGrades` view in Problem (c) solve the following queries. Note that the output of each of these queries is a nested relation.  
(i) Find each pair  $(c, s)$  where `c` is the `cno` of a course and `s` is the `sid` of a student who received an 'A' in course `c`. The type of your answer relation should be  $(cno:text, sid:text)$ .