

B561 Assignment 7
Spring 22
Indexes & Indexing
Due: 21st April 2022, 11:59pm EST

This assignment covers the following topics -

- Indexes & Indexing
- B⁺ Trees
- Hashing

In addition, you should read the following sections in Chapter 8, 14 and 15 in the textbook *Database Systems The Complete Book* by Garcia-Molina, Ullman, and Widom:

- Section 8.4.2: Some Useful Indexes
- Section 14.1: Index-structure Basics
- Section 14.2: B-Trees
- Section 14.3: Hashing and Extensible Hashing

Note that the section numbers may change depending on the version of the book you have.

Problems 4 and 7 - 20 points. All other problems - 10 points.

Submit all your answers in `assignment7.pdf`

1 Indexes & Indexing

Discussion PostgreSQL permits the creation of a variety of indexes on tables. We will review such **index creation** and examine their impact on data lookup and query processing. For more details, see the PostgreSQL manual:

<https://www.postgresql.org/docs/13/indexes.html>

Example 1 *The following SQL statements create indexes on columns or combinations of columns of the `personSkill` relation.¹ Notice that there are*

$$2^{\text{arity}(\text{personSkill})} - 1 = 2^2 - 1 = 3$$

such possible indexes.

```
create index pid_index on personSkill (pid);           --
index on pid attribute
create index skill_index on personSkill (skill);       --
index on skill attribute
create index pid_skill_index on personSkill (pid,skill); --
index (pid, skill)
```

Example 2 *It is possible to declare the type of index: `btree` or `hash`. When no index type is specified, the default is `btree`. If instead of a `Btree`, a hash index is desired, then it is necessary to specify a `using hash` qualifier:*

```
create index pid_hash on personSkill using hash (pid);
-- hash index on pid attribute
```

Example 3 *It is possible to create an index on a relation based on a scalar expression or a function defined over the attributes of that relation. Consider the following (immutable) function which computes the number of skills of a person:*

```
create or replace function numberOfSkills(p integer) returns integer as
$$
    select count(1)::int
    from   personSkill
    where  pid = p;
$$ language SQL immutable;
```

¹Incidentally, when a primary key is declared when a relation is created, PostgreSQL will create a btree index on this key for the relation.

Then the following is an index defined on the `numberOfSkills` values of persons:

```
create index numberOfSkills_index on personSkill (  
numberOfSkills(pid));
```

Such an index is useful for queries that use this function such as

```
select pid, skill from personSkill where  
numberOfSkills(pid) > 2;
```

1. Consider the relation **Company** populated with thousands of records. For the purposes of this question, let us assume that the attribute **cname** is **not** a primary-key, and has many duplicate values. On the other hand the attribute **headquarter** has relatively fewer duplicates. Which attribute of **Company** is better suited to be indexed? Explain your answer in the context of loading data into memory once the index has been created and a query is executed on the relation.
2. Let us assume a relation **Transaction**(**tid**: int, **timestamp**: date, **amount**: float) in a banking application. Assume that this relation undergoes 10000 modifications(inserts/deletes/updates) every second. What advantages or disadvantages would creating an index on any attribute have on this table? (Note that **tid** is not a primary key here.)
3. Consider the following parameters:


block size	=	4096 bytes
block-address size	=	12 bytes
block access time (I/O operation)	=	15 μs
record size	=	150 bytes
record primary key size	=	10 bytes

Assume that there is a B^+ -tree, adhering to these parameters, that indexes 2.5 billion records on their primary key values. Provide answers to the following problems and show the intermediate computations leading to your answers.

- (a) Specify (in milliseconds) the minimum time to determine if there is a record with key k in the B^+ -tree. (In this problem, you can not assume that a key value that appears in a non-leaf node has a corresponding record with that key value.)
- (b) Specify (in milliseconds) the maximum time to insert a record with key k in the B^+ -tree assuming that this record was not already in the data file.
- (c) How large must main memory be to hold the first two levels of the B^+ -tree?

4. Consider the following B^+ -tree of **order 2** that holds records with keys 4, 2, and 7.²

Strategy for splitting leaf nodes: when a leaf node n needs to be split into two nodes n_1 and n_2 (where n_1 will point to n_2), then use the rule that an even number of keys in n is moved into n_1 and an odd number of keys is moved into n_2 . So if n becomes conceptually $\boxed{k_1|k_2|k_3}$ then n_1 should be $\boxed{k_1|k_2}$ and n_2 should be $\boxed{k_3|}$ and $n_1 \rightarrow n_2$.



- Show the contents of your B^+ -tree after inserting records with keys $8 \rightarrow 9 \rightarrow 3 \rightarrow 6 \rightarrow 12$.
- Show the contents of your B^+ -tree after deleting records with keys $9 \rightarrow 7 \rightarrow 3 \rightarrow 2$.

Note: You need to show the structure of the B^+ -tree after an insert/delete operation for **each** key in the order presented in the question. Credit will not be given for directly presenting the final structure of the tree.

²Recall that (a) an internal node of a B^+ -tree of **order 2** can have either 1 or 2 keys values, and 2 or 3 sub-trees, and (b) a leaf node can have either 1 or 2 key values.

Generating data to test Indexes

For the problems below, create appropriate indexes for the **Person** (pid, pname) and **worksFor**(pid, cname, salary) relations. You need to test the queries below after populating the table with:

- (a) 100 records
- (b) 1000 records
- (c) 10000 records

Generate data using the `insertNewRecords(recSize int)` function present in `dataGen.sql`. This function will populate **Person** and **worksFor** relations with `recSize` randomized records.

For example, after creating the tables **Company**, **Person** and **worksFor**, call `insertNewRecords(100)` to create a set of 100 randomized records in the **person** and **worksFor** relations.

Note: Your task is to illustrate the speedup in these queries. One suggestion is to create a table and list your `create index` query next to the table as follows :-

Record size	Exec. time (without index)	Exec. time (with index)
100		
1000		
10000		

Appropriate Index: `create index ...`

It is recommended that you use the **pgAdmin** GUI to perform these tests as it makes things easy.

Use the `vacuum` command to clear your cache after running an experiment and use the `explain analyze` command to get the execution time of the query.

For example, `explain analyze select * from Person.`

<https://www.postgresql.org/docs/current/sql-explain.html>

<https://www.postgresql.org/docs/current/sql-vacuum.html>

5. Create an appropriate index on the **worksFor** relation that speeds up the range query

```
select pid, cname
from   worksFor
where  salary between s1 and s2;
```


Here **s1** and **s2** are two salaries with $s1 < s2$.

- (a) **Illustrate this speedup** by finding the execution times for this query for various sizes of the **worksFor** relation.
 - (b) Theoretically speaking, which index (**btree** or **hash**) is suited better for this query? Explain your reasoning. What happens to the query execution time across different range sizes across **s1** and **s2**?
6. Create indexes on the **worksFor** and **Person** relation that speed up the multiple conditions query

```
select pid, pname
from   Person
where  pid in (select pid from worksFor where cname = c)
```

Here **c** is the cname of a company. **Illustrate this speedup** by finding the execution times for this query for various sizes of the **worksFor** relation.

7. Consider an extensible hashing data structure wherein (1) the initial global depth is set at 1 and (2) all directory pointers point to the same **empty** bucket which has local depth 0. So the hashing structure looks like this:



Assume that a bucket can hold at most two records.

- (a) Show the state of the hash data structure after each of the following insert sequences:³
- i. records with keys 1 and 3.
 - ii. records with keys 6 and 9.
 - iii. records with keys 4 and 5.
 - iv. records with keys 7 and 11.
- (b) Starting from the answer you obtained for Question 7a, show the state of the hash data structure after each of the following delete sequences:
- i. records with keys 4 and 11.
 - ii. records with keys 9 and 7.
 - iii. records with keys 1 and 5.

As in the text book, the bit strings are interpreted starting from their left-most bit and continuing to the right subsequent bits.

8. In practice, several hash functions don't operate as well as expected. Briefly explain the disadvantages of **linear hashing**. Consider a **linear hash function** for integer keys i :

$$h(i) = i^2 \bmod(B)$$

where B is a positive integer that represents the number of buckets and **mod** is the modulo function.

- (a) Evaluate $h(i)$ for $B = 10$.
- (b) Evaluate $h(i)$ for $B = 16$.
- (c) What would be the best value of B for $h(i)$?

In your answers, use the intuition you gained from explaining the disadvantages of linear hashing, and use diagrams (if necessary) to support your evaluation.

³For Problem 7 - you should interpret the key values as bit strings of length 4. So for example, key value 7 is represented as the bit string 0111 and key value 2 is represented as the bit string 0010.